

Di Hu

dihu@ruc.edu.cn | [personal homepage](#) | [lab homepage](#)

WORKING EXPERIENCE

Tenure-track Associate Professor <i>Gaoling School of Artificial Intelligence, Renmin University of China</i>	Jul. 2023 – Present
Tenure-track Assistant Professor <i>Gaoling School of Artificial Intelligence, Renmin University of China</i>	Aug. 2020 – Jul. 2023
Research Scientist <i>Baidu Research</i>	Jul. 2019 – Aug. 2020

EDUCATION

Northwestern Polytechnical University <i>Ph.D in Computer Science and Technologys</i> <i>Advisor: Feiping Nie, Xuelong Li</i> <i>Thesis: Research on Machine Multimodal Perception</i>	2014 – 2019
Honor College, Northwestern Polytechnical University <i>Bachelor in Computer Science and Technology</i>	2010 – 2014

RESEARCH INTEREST

Interested in how to understand and interact with the environment via the natural multimodal messages, *e.g.*, *sound, vision and touch*. I'm strongly convinced that the pervasive multimodal messages can provide sufficient information for perceiving, interacting, learning and understanding environment, even the agent itself, which promisingly makes multimodal learning become one of the key to achieve machine general intelligence.

DISTINCTION

BAAI Scholar	2025
MSRA StarTrack Scholar	2025
WuWenJun AI Excellent Young Scientist Award	2023
The Young Elite Scientists Sponsorship Program	2022
SHAANXI Outstanding Doctoral Dissertation Award	2021
CAAI Outstanding Doctoral Dissertation Award	2020
ACM Xi'an Doctoral Dissertation Award	2019
Baidu AIDU Recruitment Program	2019
CVPR Doctoral Consortium	2019
CSC Scholarship to CMU as a Visiting Scholar	2018
National Scholarship	2017, 2018
RoboCup China Open <i>The First Prize, Service Robot@Home</i>	2014

PUBLICATIONS

Conference Paper	(†: Equal Contribution, *: Corresponding Author)
52. Yake Wei, Yu Miao, Dongzhan Zhou, and Di Hu* . <i>MokA: Multimodal Low-Rank Adaptation for MLLMs</i> . In <i>Advances in Neural Information Processing Systems (NeurIPS)</i> , 2025. Oral Presentation	
51. Wenke Xia, Yichu Yang, Hongtao Wu, Xiao Ma, Tao Kong, and Di Hu* . Robotic Policy Learning via Human-assisted Action Preference Optimization. In <i>Advances in Neural Information Processing Systems (NeurIPS)</i> , 2025.	
50. Zequn Yang, Hongfa Wang, and Di Hu* . Efficient Quantification of Multimodal Interaction at Sample Level. In <i>Proceedings of the International Conference on Machine Learning (ICML)</i> , 2025.	

49. Haotian Ni, Yake Wei, Hang Liu, Gong Chen, Chong Peng, Hao Lin, and **Di Hu***. Reviving the Cooperation Dynamics in Multimodal Transformer. *In Proceedings of the International Conference on Machine Learning (ICML)*, 2025.
48. Wenke Xia, Ruoxuan Feng, Dong Wang, and **Di Hu***. Phoenix: A Motion-based Self-Reflection Framework for Fine-grained Robotic Action Correction. *In Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, 2025.
47. Henghui Du, Guangyao Li, Chang Zhou, Chunjie Zhang, Alan Zhao, and **Di Hu***. Crab: A Unified Audio-Visual Scene Understanding Model with Explicit Cooperation. *In Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, 2025.
46. Chengxiang Huang[†], Yake Wei[†], Zequn Yang, and **Di Hu***. Adaptive Unimodal Regulation for Balanced Multimodal Information Acquisition. *In Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, 2025.
45. Ruotian Peng[†], Haiying He[†], Yake Wei, Yandong Wen, and **Di Hu***. Patch Matters: Training-free Fine-grained Image Caption Enhancement via Local Perception. *In Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, 2025.
44. Ruoxuan Feng, Zhen Tian, Qiushi Peng, Jiaxin Mao, Xin Zhao, **Di Hu*** and Changwang Zhang. MGIPF: Multi-Granularity Interest Prediction Framework for Personalized Recommendation. *In Proceedings of the International ACM SIGIR Conference on Research and Development in Information Retrieval (SIGIR)*, 2025.
43. Ruoxuan Feng, Jiangyu Hu, Wenke Xia, TianciGao, Ao Shen, Yuhao Sun, Bin Fang*, **Di Hu***. AnyTouch: Learning Unified Static-Dynamic Representation across Multiple Visuo-tactile Sensors. *In Proceedings of the International Conference on Learning Representations (ICLR)*, 2025.
42. Ruoxuan Feng, **Di Hu***, Wenke Ma, and Xuelong Li. Play to the Score: Stage-Guided Dynamic Multi-Sensory Fusion for Robotic Manipulation. *In Conference on Robot Learning (CoRL)*, 2024. **Oral Presentation**
41. Jingxian Lu, Wenke Xia, Dong Wang, Zhigang Wang, Bin Zhao, **Di Hu***, and Xuelong Li. KOI: Accelerating Online Imitation Learning via Hybrid Key-state Guidance. *In Conference on Robot Learning (CoRL)*, 2024.
40. Yake Wei, Siwei Li, Ruoxuan Feng, and **Di Hu***. Diagnosing and Re-learning for Balanced Multimodal Learning. *In Proceedings of the European Conference on Computer Vision (ECCV)*, 2024.
39. Juncheng Ma, Peiwen Sun, Yaoting Wang, and **Di Hu***. Stepping Stones: A Progressive Training Strategy for Audio-Visual Semantic Segmentation. *In Proceedings of the European Conference on Computer Vision (ECCV)*, 2024.
38. Yaoting Wang[†], Peiwen Sun[†], Dongzhan Zhou, Guangyao Li, Honggang Zhang, and **Di Hu***. Ref-AVS: Refer and Segment Objects in Audio-Visual Scenes. *In Proceedings of the European Conference on Computer Vision (ECCV)*, 2024.
37. Yaoting Wang[†], Peiwen Sun[†], Yuanchao Li, Honggang Zhang, and **Di Hu***. Can Textual Semantics Mitigate Sounding Object Segmentation Preference? *In Proceedings of the European Conference on Computer Vision (ECCV)*, 2024.
36. Guangyao Li, HenghuiDu, and **Di Hu***. Boosting Audio Visual Question Answering via Key Semantic-Aware Cues. *In Proceedings of the ACM Conference on Multimedia (ACMMM)*, 2024.
35. Peiwen Sun, Honggang Zhang, and **Di Hu***. Unveiling and Mitigating Bias in Audio Visual Segmentation. *In Proceedings of the ACM Conference on Multimedia (ACMMM)*, 2024. **Oral Presentation**
34. Xincheng Pang[†], Wenke Xia[†], Zhigang Wang, Bin Zhao, **Di Hu***, Dong Wang*, and Xuelong Li. Depth Helps: Improving Pre-trained RGB-based Policy with Depth Information Injection. *In Proceedings of the International Conference on Intelligent Robots and Systems (IROS)*, 2024.
33. Jia Zeng, Qingwen Bu, Bangjun Wang, Wenke Xia, Li Chen, Hao Dong, Haoming Song, Dong Wang, **Di Hu**, Ping Luo, Heming Cui, Bin Zhao, Xuelong Li, Yu Qiao, and Hongyang Li. Learning Manipulation by Predicting Interaction. *In Proceedings of Robotics: Science and Systems Conference (RSS)*, 2024.
32. Yake Wei and **Di Hu***. MMPareto: Innocent Uni-modal Assistance for Enhanced Multi-modal Learning. *In Proceedings of the International Conference on Machine Learning (ICML)*, 2024.
31. Yake Wei, Ruoxuan Feng, Ziheng Wang, and **Di Hu***. Enhancing Multi-modal Cooperation via Fine-grained Modality Valuation. *In Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, 2024.
30. Wenke Xia, Dong Wang, Xincheng Pang, Zhigang Wang, Bin Zhao, **Di Hu***, and Xuelong Li. Kinematic-aware Prompting for Generalizable Articulated Object Manipulation with LLMs. *In Proceedings of the International Conference on Robotics and Automation (ICRA)*, 2024.

29. Zequn Yang, Yake Wei, Ce Liang, and **Di Hu***. Quantifying and Enhancing Multi-modal Robustness with Modality Preference. In *Proceedings of the International Conference on Learning Representations (ICLR)*, 2024.
28. Yaoting Wang[†], Weisong Liu[†], Guangyao Li, Jian Ding, **Di Hu***, and Xi Li. Prompting Segmentation with Sound is Generalizable Audio-Visual Source Localizer. In *Proceedings of the AAAI Conference on Artificial Intelligence (AAAI)*, 2024.
27. Tao Wu, Xuewei Li, Zhongang Qi, **Di Hu**, Xintao Wang, Ying Shan, and Xi Li. SphereDiffusion: Spherical Geometry-aware Distortion Resilient Diffusion Model. In *Proceedings of the AAAI Conference on Artificial Intelligence (AAAI)*, 2024.
26. Guangyao Li, Wenxuan Hou, and **Di Hu***. Progressive Spatio-temporal Perception for Audio-Visual Question Answering. In *Proceedings of the ACM Conference on Multimedia (ACMMM)*, 2023.
25. Hongpeng Lin[†], Ludan Ruan[†], Wenke Xia[†], Peiyu Liu, Jingyuan Wen, Yixin Xu, **Di Hu**, Ruihua Song, Wayne Xin Zhao, Qin Jin, and Zhiwu Lu. TikTok: A Video-Based Dialogue Dataset for Multi-Modal Chitchat in Real World. In *Proceedings of the ACM Conference on Multimedia (ACMMM)*, 2023.
24. Andong Deng, Xingjian Li, **Di Hu***, Tianyang Wang, Haoyi Xiong, Chengzhong Xu. Towards Inadequately Pre-trained Models in Transfer Learning. In *Proceedings of the IEEE Conference on Computer Vision (ICCV)*, 2023.
23. Guangyao Li, Yixin Xu, and **Di Hu***. Multi-Scale Attention for Audio Question Answering. *Interspeech*, 2023.
22. Wenke Xia, Xingjian Li, Andong Deng, Haoyi Xiong, Dejing Dou, and **Di Hu***. Robust Cross-Modal Knowledge Distillation for Unconstrained Videos. *IEEE International Conference on Multimedia and Expo (ICME)*, 2023.
21. Ruize Xu, Ruoxuan Feng, Shi-xiong Zhang, and **Di Hu***. MMCosine: Multi-Modal Cosine Loss Towards Balanced Audio-Visual Fine-Grained Learning. *The International Conference on Acoustics, Speech, & Signal Processing (ICASSP)*, 2023.
20. Xinchu Zhou, Dongzhan Zhou, **Di Hu**, Hang Zhou, and Wanli Ouyang. Exploiting Visual Context Semantics for Sound Source Localization. *IEEE/CVF Winter Conference on Applications of Computer Vision (WACV)*, 2022.
19. Xinchu Zhou, Dongzhan Zhou, Wanli Ouyang, Hang Zhou, and **Di Hu**. SeCo: Separating Unknown Musical Visual Sounds with Consistency Guidance. *IEEE/CVF Winter Conference on Applications of Computer Vision (WACV)*, 2022.
18. Xiaokang Peng[†], Yake Wei[†], Andong Deng, Dong Wang, and **Di Hu***. Balanced Multimodal Learning via On-the-fly Gradient Modulation. In *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, 2022. **Oral Presentation**
17. Guangyao Li[†], Yake Wei[†], Yapeng Tian[†], Chenliang Xu, Ji-Rong Wen, and **Di Hu***. Learning to Answer Questions in Dynamic Audio-Visual Scenarios. In *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, 2022. **Oral Presentation**
16. Xian Liu, Rui Qian, Hang Zhou, **Di Hu**, Weiyao Lin, Ziwei Liu, Bolei Zhou, and Xiaowei Zhou. Visual Sound Localization in-the-Wild by Cross-Modal Interference Erasing. In *Proceedings of the AAAI Conference on Artificial Intelligence (AAAI)*, 2022.
15. Dongzhan Zhou, Xinchu Zhou, **Di Hu***, Hang Zhou, Lei Bai, Ziwei Liu, and Wanli Ouyang. SepFusion: Finding Optimal Fusion Structures for Visual Sound Separation. In *Proceedings of the AAAI Conference on Artificial Intelligence (AAAI)*, 2022.
14. Zechen Bai, Zhigang Wang, Jian Wang, **Di Hu***, and Errui Ding*. Unsupervised Multi-Source Domain Adaptation for Person Re-Identification. In *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, 2021. **Oral Presentation**
13. Yapeng Tian, **Di Hu***, and Chenliang Xu*. Cyclic Co-Learning of Sounding Object Visual Grounding and Sound Separation. In *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, 2021.
12. Dong Wang, **Di Hu***, Xingjian Li, and Dejing Dou. Temporal Relational Modeling with Self-Supervision for Action Segmentation. In *Proceedings of the AAAI Conference on Artificial Intelligence (AAAI)*, 2021.
11. **Di Hu**, Rui Qian, Minyue Jiang, Xiao Tan, Shilei Wen, Errui Ding, Weiyao Lin, and Dejing Dou. Discriminative Sounding Objects Localization via Self-supervised Audiovisual Matching. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2020.
10. **Di Hu**, Xuhong Li, Lichao Mou, Pu Jin, Dong Chen, Liping Jing, Xiaoxiang Zhu, and Dejing Dou. Cross-Task Transfer for Geotagged Audiovisual Aerial Scene Recognition. In *Proceedings of the European Conference on Computer Vision (ECCV)*, 2020.
9. Rui Qian, **Di Hu**, Heinrich Dinkel, Mengyue Wu, Ning Xu, and Weiyao Lin. Multiple Sound Sources Localization from Coarse to Fine. In *Proceedings of the European Conference on Computer Vision (ECCV)*, 2020.

8. **Di Hu**, Dong Wang, Xuelong Li, Feiping Nie, and Qi Wang. Listen to the Image. *In Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, 2019.
7. **Di Hu**, Feiping Nie, and Xuelong Li. Deep Multimodal Clustering for Unsupervised Audiovisual Learning. *In Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, 2019.
6. **Di Hu**, Chengze Wang, Feiping Nie, and Xuelong Li. Dense Multimodal Fusion for Hierarchically Joint Representation. *In Proceedings of the IEEE Conference on Acoustics, Speech and Signal Processing (ICASSP)*, 2019.
5. Xuelong Li, **Di Hu**, and Feiping Nie. Large Graph Hashing with Spectral Rotation. *In Proceedings of the AAAI Conference on Artificial Intelligence (AAAI)*, 2017.
4. Xuelong Li, **Di Hu**, and Feiping Nie. Deep Binary Reconstruction for Cross-modal Hashing. *In Proceedings of the ACM Conference on Multimedia (ACMMM)*, 2017.
3. Xuelong Li, **Di Hu**, and Xiaoqiang Lu. Image2song: Song Retrieval via Bridging Image Content and Lyric Words. *In Proceedings of the IEEE Conference on Computer Vision (ICCV)*, 2017.
2. **Di Hu**, Xiaoqiang Lu, and Xuelong Li. Multimodal Learning via Exploring Deep Semantic Similarity. *In Proceedings of the ACM Conference on Multimedia (ACMMM)*, 2016.
1. **Di Hu**, Xuelong Li, and Xiaoqiang Lu. Temporal Multimodal Learning in Audiovisual Speech Recognition. *In Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, 2016.

Journal Paper

11. Yake Wei, **Di Hu***, Henghui Du, and Ji-Rong Wen. On-the-fly Modulation for Balanced Multimodal Learning. *In IEEE Trans. Pattern Analysis and Machine Intelligence (TPAMI)*, 2024.
10. Wenxuan Hou†, Guangyao Li†, Yapeng Tian, and **Di Hu***. Towards Long Form Audio-visual Video Understanding. *ACM Transactions on Multimedia Computing, Communications, and Applications (TOMM)*, 2024.
9. Zequn Yang, Han Zhang, Yake Wei, Zheng Wang, Feiping Nie, and **Di Hu***. Geometric-Inspired Graph-based Incomplete Multi-view Clustering. *Pattern Recognition (PR)*, 2024.
8. Ziyun Li, Jona Otholt, Ben Dai, **Di Hu**, Christoph Meinel, and Haojin Yang. Supervised Knowledge May Hurt Novel Class Discovery Performance. *Transactions on Machine Learning Research (TMLR)*, 2023.
7. Konrad Heidler, Lichao Mou, **Di Hu**, Pu Jin, Guangyao Li, Chuang Gan, Ji-Rong Wen, Xiao-Xiang Zhu. Self-supervised Audiovisual Representation Learning for Remote Sensing Data. *In International Journal of Applied Earth Observation and Geoinformation*, 2023.
6. **Di Hu**, Zheng Wang, Feiping Nie, Rong Wang, Xuelong Li. Self-supervised Learning for Heterogeneous Audiovisual Scene Analysis. *In IEEE Trans. Multimedia (TMM)*, 2022.
5. **Di Hu**, Yake Wei, Rui Qian, Weiyao Lin, Ruihua Song, Ji-Rong Wen. Class-aware Sounding Objects Localization via Audiovisual Correspondence. *In IEEE Trans. Pattern Analysis and Machine Intelligence (TPAMI)*, 2021.
4. Sijia Yang, Haoyi Xiong, **Di Hu**, Kaibo Xu, Licheng Wang, Peizhen Zhu, Zeyi Sun. Generalising Combinatorial Discriminant Analysis through Conditioning Truncated Rayleigh Flow. *Knowledge and Information Systems (KAIS)*, 2021.
3. **Di Hu**, Feiping Nie, and Xuelong Li. Deep Linear Discriminant Analysis Hashing. *In SCIENTIA SINICA Informationis*, 2019.
2. **Di Hu**, Feiping Nie, and Xuelong Li. Discrete Spectral Hashing for Efficient Similarity Retrieval. *In IEEE Trans. Image Processing (TIP)*, 2018.
1. **Di Hu**, Feiping Nie, and Xuelong Li. Deep Binary Reconstruction for Cross-modal Hashing. *In IEEE Trans. Multimedia (TMM)*, 2018.

Workshop Paper

6. Guangyao Li, Henghui Du, and **Di Hu***. AVQA-CoT: When CoT Meets Question Answering in Audio-Visual Scenarios. *In Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition Workshop (CVPRW)*, 2024.
5. Wenke Xia†, Xu Zhao†, Xincheng Pang, Changqing Zhang, and **Di Hu***. Balanced Audiovisual Dataset for Imbalance Analysis. *In Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition Workshop (CVPRW)*, 2023.
4. **Di Hu**, Zheng Wang, Haoyi Xiong, Dong Wang, Feiping Nie, and Dejing Dou. Heterogeneous Scene Analysis via Self-supervised Audiovisual Learning. *In Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition Workshop (CVPRW)*, 2020.

3. **Di Hu***, Lichao Mou*, Qingzhong Wang*, Junyu Gao, Yuansheng Hua, Dejing Dou, and Xiaoxiang Zhu. Does Ambient Sound Help? - Audiovisual Crowd Counting. *In Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition Workshop (CVPRW)*, 2020.
2. Yapeng Tian*, **Di Hu***, and Chenliang Xu. Co-Learn Sounding Object Visual Grounding and Visually Indicated Sound Separation in A Cycle Video. *In Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition Workshop (CVPRW)*, 2020.
1. Rui Qian, **Di Hu**, Heinrich Dinkel, Mengyue Wu, Ning Xu, and Weiyao Lin. A Two-Stage Framework for Multiple Sound-Source Localization. *In Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition Workshop (CVPRW)*, 2020.

PROFESSIONAL SERVICES

Organizing Committee

CVPR Tutorial on Audio-visual Scene Understanding	2021
WACV Tutorial on Audio-visual Scene Understanding	2021
ICDM Tutorial on Automated Deep Learning: Theory, Algorithms, Platforms, and Applications	2019

Senior Program Committee / Area Chair

International Conference on Learning Representations (ICLR)	2026
The AAAI Conference on Artificial Intelligence (AAAI)	2023-2025
The International Joint Conference on Artificial Intelligence (IJCAI)	2023-2025

Program Committee

IEEE Conference on Computer Vision and Pattern Recognition (CVPR)	2018, 2020-2025
IEEE International Conference on Computer Vision (ICCV)	2019, 2021, 2025
European Conference on Computer Vision (ECCV)	2020, 2022
The AAAI Conference on Artificial Intelligence (AAAI)	2018, 2020-2022
International Conference in Learning Representations (ICLR)	2021-2024
Neural Information Processing Systems (NeurIPS)	2020-2024
The International Conference on Machine Learning (ICML)	2021-2025
Asian Conference on Computer Vision (ACCV)	2018, 2020
IEEE Winter Conference on Applications of Computer Vision (WACV)	2021

Journal Reviewer

IEEE Transactions on Pattern Analysis and Machine Intelligence (TPAMI)
IEEE Transactions on Neural Networks and Learning Systems (TNNLS)
IEEE Transactions on Image Processing (TIP)
IEEE Transactions on Circuits and Systems for Video Technology (TCSVT)
IEEE Transactions on Multimedia (TMM)
IEEE Transactions on Knowledge and Data Engineering (TKDE)
ACM Transactions on Intelligent Systems and Technology (TIST)